



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

Parallelizing 2D Ising model

Different implementations and benchmarks

W. Conte

February 19, 2026



Table of Contents

1. Introduction

- Physics
- Hardware

2. Implementation

- General structure
- Serial implementation
- Parallel implementations

3. Benchmarks

- Physics correctness
- Performances



Table of Contents

1. Introduction

- Physics
- Hardware

2. Implementation

- General structure
- Serial implementation
- Parallel implementations

3. Benchmarks

- Physics correctness
- Performances

Ising Model (1)



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

Ising model is a statistical mechanics framework used to describe ferromagnetic systems.

1. The Hamiltonian (Energy)

The energy of a configuration σ is defined by:

$$\mathcal{H}(\sigma) = -J \sum_{\langle i,j \rangle} \sigma_i \sigma_j - h \sum_i \sigma_i$$

- $J > 0$: Ferromagnetic coupling constant (neighbors prefer alignment).
- $\langle i,j \rangle$: Sum over nearest neighbors (top, bottom, left, right).
- h : External magnetic field (set to 0 for spontaneous magnetization tests).



2. Order Parameter (Magnetization)

The macroscopic magnetization acts as the order parameter for the phase transition:

$$m = \frac{1}{N} \sum_{i=1}^N \sigma_i \quad \in [-1, 1]$$

- **High T :** $m \approx 0$ (Paramagnetic / Disordered Phase).
- **Low T :** $|m| \rightarrow 1$ (Ferromagnetic / Ordered Phase).

The Metropolis-Hastings Algorithm



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

Core Idea: A Markov Chain Monte Carlo (MCMC) method to sample configurations according to the Boltzmann distribution $P(\sigma) \propto e^{-\beta \mathcal{H}(\sigma)}$.

Why use it?

- Calculating the partition function Z is analytically intractable for large N .
- We need a way to simulate thermal fluctuations and reach equilibrium efficiently.

How it works (The Step):

1. Pick a random spin σ_i and propose a flip: $\sigma'_i = -\sigma_i$.
2. Calculate energy change:

$$\Delta E = 2\sigma_i \left(J \sum_{j \in nn(i)} \sigma_j + h \right)$$

3. Acceptance Criterion:

- If $\Delta E \leq 0$: **Accept**
- If $\Delta E > 0$: Accept with probability $p = e^{-\beta \Delta E}$.

Hardware specifications: the host (CPU)



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

Processor: Intel® Core™ i5-13420H

Hybrid Architecture Configuration:

Total: 8 Cores / 12 Threads

1. Performance Cores (P-Cores)

- **Count:** 4 Physical Cores (8 Threads via Hyper-Threading).
- **Role:** High-performance tasks, low latency, heavy single-thread workload.

2. Efficiency Cores (E-Cores)

- **Count:** 4 Physical Cores (4 Threads, NO Hyper-Threading).
- **Role:** Background tasks
- **Characteristics:** Lower clock speeds, optimized for power efficiency.

Hardware specifications: the device (GPU)



Device: NVIDIA GeForce RTX 4050 Laptop GPU

Specs:

- **Architecture:** Ada Lovelace
- **CUDA Cores:** 2560
- **VRAM:** 6 GB
- **Memory Bandwidth:** 192 GB/s
- **SM count:** 20 (128 CUDA cores each)
- **Max threads / SM:** 1536

The "Ada" Advantage:

- Unlike previous generations (Ampere), Ada Lovelace features a massive increase in on-chip memory.
- **L2 Cache Size:** 24 MB



Table of Contents

1. Introduction

- Physics
- Hardware

2. Implementation

- General structure
- Serial implementation
- Parallel implementations

3. Benchmarks

- Physics correctness
- Performances

Class Structure and Wrapper



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

```
● ● ●  
// ...  
  
enum class Mode {serial, openMP, cuda_global, cuda_shared};  
  
class IsingModel2d{  
public:  
  
    // constructor  
    IsingModel2d(int L, double T, double J, double h, unsigned int seed);  
    // destructor  
    ~IsingModel2d();  
  
    // default cuda block size (only one side)  
    int cuda_block_size = 16;  
  
    // ...  
  
private:  
  
    int L;  
    int row_stride; // is equal to L + 2, because it includes the padding  
    double T;  
    double beta; // = 1/T  
    double J = 1; // interaction term  
    double h = 0; // magnetic field  
    // lookup table: this will be essentially a probability grid to  
    // avoid computing exp every time  
    float lookup_probs[2][5];  
  
    // CUDA Device Pointers  
    int* d_lattice = nullptr;  
    float* d_lookup_probs = nullptr;  
    // Pointer to a curandState which is on the device  
    void* d_states = nullptr; // void* to hide curandState to g++  
  
    // lattice is already flattened in 1D vector for simplicity on next phases  
    std::vector<int> lattice;  
  
    // random number generators  
    std::mt19937 serial_rng;  
    std::vector<std::mt19937> omp_rngs; // one number per thread to avoid race conditions  
    unsigned int n_seed;  
  
    // ...
```

- enum class for “evolution” modes (Serial, OpenMP, CUDA)
- **Public Methods** for physical observables, updates, and memory management
- **Encapsulation:** Private members handle internal logic via CUDA “wrapper functions”
- **Interoperability:** Python wrapper implemented using PyBind11

Constructor - generalities



```
IsingModel2d::IsingModel2d(int L, double T, double J, double h, unsigned int seed) : L(L), row_stride(L + 2), T(T),
beta(1.0/T), J(J), h(h), serial_rng(seed),
m_seed(seed) {
    // initialize lattice with total size including padding
    lattice.resize(row_stride * row_stride);

    // random number generator for spins
    std::uniform_int_distribution<int> uni_dist(0, 1);

    // initialize core spins to +1 or -1
    for (int i = 1; i <= L; i++) {
        for (int j = 1; j <= L; j++) {
            lattice[i * row_stride + j] = 2 * uni_dist(serial_rng) - 1;
        }
    }
}
```

- Lattice mapped to 1D array (row-major) → ensures contiguous memory layout

Lookup probability table



```
// ...  
  
// Compute lookup probabilities to determine spin flips during Metropolis updates.  
// In this way we compute only once and we do only "read" operation in the following,  
// which are more efficient than compute an exponential function at each step  
for (int s_idx = 0; s_idx < 2; ++s_idx) {  
    // assign to index 0 the values of spin -1 and  
    // to index 1 the values of spin +1  
    double s = (s_idx == 0) ? -1.0 : 1.0;  
  
    // for a given spin value, store the probability of spin flip given the neighbors  
    for (int i = 0; i < 5; ++i) {  
        // Index mapping: index 0 is relative to value of neighbors -4 (i.e., all negative)  
        int physical_sum = (i * 2) - 4; // possible sum of neighbors is -4,-2,0,2,4  
        // Globally we have this index mapping: (sum_neighbors --> index_lookup_table)  
        // -4 --> 0 | -2 --> 1 | 0 --> 2 | 2 --> 3 | 4 --> 3  
  
        // compute delta E  
        double delta_E = 2.0 * s * (J * physical_sum + h);  
  
        // assign to each position in the lookup prob the correspondent acceptance probability  
        lookup_probs[s_idx][i] = (delta_E > 0) ? std::exp(-delta_E * beta) : 1.0;  
    }  
}
```

- Pre-computation of Boltzmann weights for all possible ΔE configurations.
- Replaces expensive exponential calls ($e^{-\beta \Delta E}$) with simple read-only memory accesses during the simulation.

Periodic Boundary Conditions - CPU



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

Implementation (ternary operator):

```
int i_up = (i == 1) ? L : i - 1;
```

"If at top edge, go to bottom; else, go up."

Why?

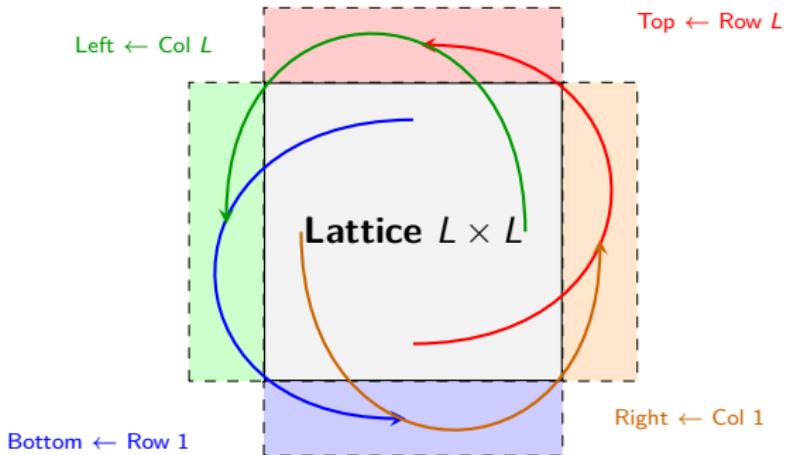
- **Speed:** Calculating an index is faster than fetching data from RAM (like using padding layers).
- **Optimization:** Branch prediction unit.

Periodic Boundary Conditions - GPU



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

- Halo Cells (Padding): the lattice is enlarged to $(L + 2) \times (L + 2)$.
- Edge elements are copied into the padding rows/columns.



Why Padding for GPU?



```
***** Sync padding *****
__global__ void sync_padding_kernel(int* lattice, int L, int row_stride) {
    int k = blockIdx.x * blockDim.x + threadIdx.x + 1;
    if (k <= L) {
        // Copy top/bottom rows
        lattice[0 * row_stride + k] = lattice[L * row_stride + k];
        lattice[(L + 1) * row_stride + k] = lattice[1 * row_stride + k];
        // Copy left/right columns
        lattice[k * row_stride + 0] = lattice[k * row_stride + L];
        lattice[k * row_stride + (L + 1)] = lattice[k * row_stride + 1];
    }
}
```

- **Memory Coalescing:** threads read contiguous memory addresses.
- **Eliminates Branch Divergence:** avoids if-else checks for boundaries inside the kernel. All threads execute the exact same instructions (SIMT)

Energy & Magnetization



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

```
double IsingModel2d::magnetization(Mode mode) {
    double m = 0.0;

    // Serial version
    if (mode == Mode::serial) {
        for (int i = 1; i <= L; i++) {
            for (int j = 1; j <= L; j++) {
                m += lattice[i * row_stride + j];
            }
        }
    }

    //...((OpenMP implementation with Reduction - not used)

    else if (mode == Mode::cuda_global || mode == Mode::cuda_shared){
        // this copy_to_host() is the dominant cost for large lattices
        copy_to_host();
        // compute magnetization on CPU
        return magnetization(Mode::serial);
    }

    return m / (double)(L * L);
}
```

Magnetization per spin

```
double IsingModel2d::energy(Mode mode) {
    double E = 0.0;

    // Serial version
    if (mode == Mode::serial){
        for (int i = 1; i <= L; i++) {
            for (int j = 1; j <= L; j++) {
                int array_index = i * row_stride + j;

                // boundary conditions with ternary operator
                int i_up = (i == 1) ? L - 1;
                int i_down = (i == L) ? 1 : i + 1;
                int j_left = (j == 1) ? L - 1 : i - 1;
                int j_right = (j == L) ? 1 : i + 1;

                int neighbors = lattice[i_up * row_stride + j] // up
                    lattice[i_down * row_stride + j] // down
                    lattice[i * row_stride + j].left // left
                    lattice[i * row_stride + j].right; // right

                // Interaction energy (halved for double counting) + field energy
                E += -0.5 * j * lattice[array_index] * neighbors + h * lattice[array_index];
            }
        }
    }

    // ... Parallel OpenMP version - NOT USED ...
    // Parallel CUDA version
    else if (mode == Mode::cuda_global || mode == Mode::cuda_shared){
        // this copy_to_host() is the dominant cost for large lattices.
        copy_to_host();

        // delegate the calculation to the CPU implementation
        return energy(Mode::serial);
    }

    return E;
}
```

Energy

SERIAL IMPLEMENTATION

Metropolis criterion



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

```
● ● ●

void IsingModel2d::Metropolis_update(int i, int j, std::mt19937& rng) {
    // Identify the 1D index of the spin we are trying to flip
    int current_idx = i * row_stride + j;

    // Identify Neighbor Indices using Periodic Boundary Conditions (PBC).
    // check if we are on a boundary (i or L) and "wrap around" explicitly.

    // if i is 1 (top row), the neighbor above is L (bottom row). Else i-1.
    int i_up = (i == 1) ? L : i - 1;
    // if i is L (bottom row), the neighbor below is 1 (top row). Else i+1.
    int i_down = (i == L) ? 1 : i + 1;
    // if j is 1 (left col), the neighbor left is L (right col). Else j-1.
    int j_left = (j == 1) ? L : j - 1;
    // if j is L (right col), the neighbor right is 1 (left col). Else j+1.
    int j_right = (j == L) ? 1 : j + 1;
    // Sum the neighbors by reading directly from the real lattice locations.
    int neighbors = lattice[i_up * row_stride + j] +
                    lattice[i_down * row_stride + j] +
                    lattice[i * row_stride + j_left] +
                    lattice[i * row_stride + j_right];

    // Calculate Energy Change (Delta E) via Lookup Table
    // map the current spin (-1 or 1) to a table index (0 or 1)
    int spin_val = lattice[current_idx];
    int spin_idx = (spin_val < 0) ? 0 : 1;

    // map the neighbor sum (-4, -2, 0, 2, 4) to a table index (0, 1, 2, 3, 4)
    int sum_idx = (neighbors + 4) / 2;

    // use a static distribution to avoid the overhead of reconstructing
    // the object at every single function call
    static std::uniform_real_distribution<double> dist(0.0, 1.0);

    // check the probability against the pre-computed lookup table.
    if (dist(rng) < lookup_probs[spin_idx][sum_idx]) {
        lattice[current_idx] *= -1; // Flip the spin
    }
}
```

PARALLEL IMPLEMENTATIONS

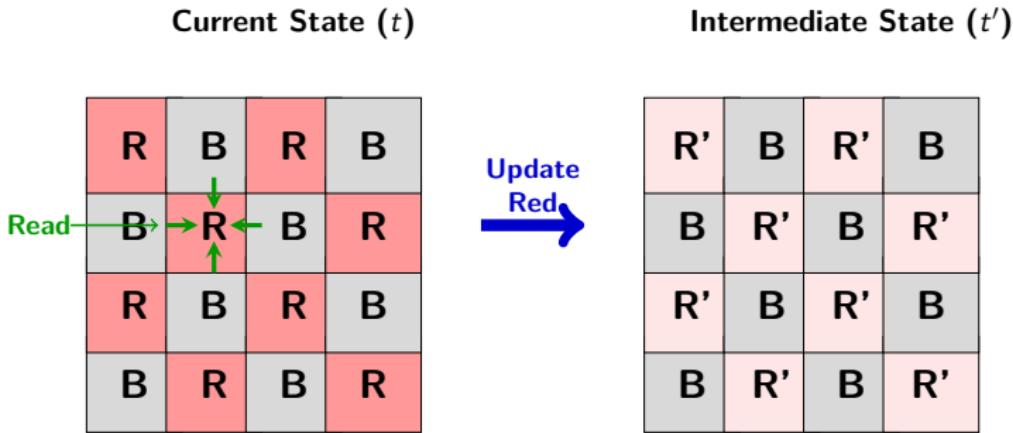
Checkerboard algorithm



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

- The update of a spin requires the state of its nearest neighbors
- **Race Condition:** If neighboring spins are updated simultaneously by different threads, the read/write order is undefined
- This leads to non-deterministic results (wrong physics) or requires costly locking mechanisms (performance kill)

Checkerboard algorithm: Step 1 (Red Update)

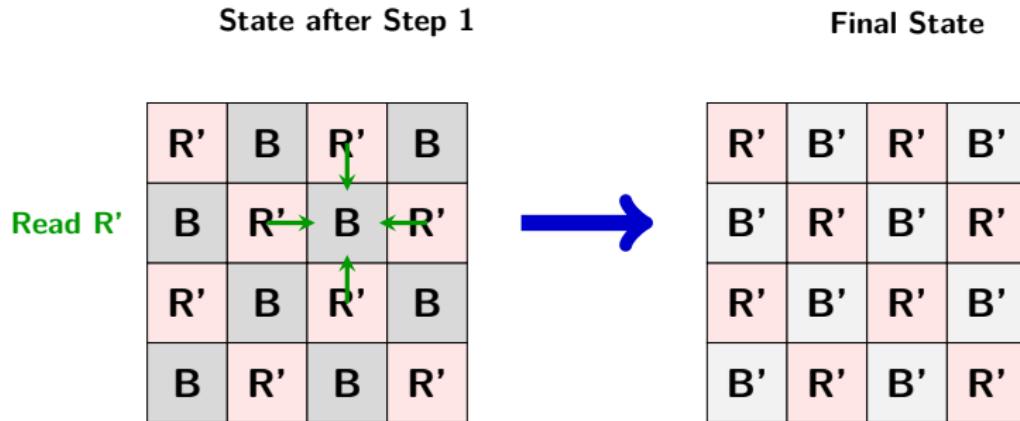


- All **Red cells** are updated simultaneously.
- They only read from Black neighbors, which are constant during this step.
- **No Race Conditions:** Threads do not write to the same memory locations concurrently.

Checkerboard Algorithm: Step 2 (Black update)



UNIVERSITÀ
DEGLI STUDI
DI PADOVA



This scheme will be applied both for OpenMP and CUDA

OpenMP - implementation (1)



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

```
● ● ●

void IsingModel2d::step_openmp(int steps) {
    //sync_padding();

    // Threads are created here and stay alive for all simulation
    #pragma omp parallel
    {
        // each thread gets its own unique ID and private RNG state
        int thread_id = omp_get_thread_num();
        std::mt19937& thread_rng = omp_rngs[thread_id];
    }
    // ...
}
```

- Start the parallel region: assign a different RNG at each thread

OpenMP - Implementation (2)



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

```
● ● ●

// ...
// time loop
for (int s = 0; s < steps; s++) {

    // even spins
    #pragma omp for collapse(2)
    for(int i = 1; i <= L; i++){
        for(int j = 1; j <= L; j++){
            // Checkerboard condition for even sites
            if ((i + j) % 2 == 0) {
                Metropolis_update(i, j, thread_rng);
            }
        }
    }

    // odd spins
    #pragma omp for collapse(2)
    for(int i = 1; i <= L; i++){
        for(int j = 1; j <= L; j++){
            // checkerboard condition for odd sites
            if ((i + j) % 2 != 0) {
                Metropolis_update(i, j, thread_rng);
            }
        }
    }

} // end of time loop
} // end of parallel region (threads are destroyed here)
}
```

1. Parallelize the nested for loop
2. Update "Red" cells
3. Parallelize the nested for loop
4. Update "Black" cells

CUDA - Implementation Strategies



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

- **Global Memory (Baseline):** Standard implementation without caching optimizations
- **Shared Memory + read only cache:**
 - **Shared Memory:** Used to store and reuse the thread block's internal lattice
 - **Read-only data cache:** Used for cached, read-only access to **block neighbors**

CUDA - Global memory kernel



```
•••  
  
__global__ void ising_step_global(int* lattice, int L, int row_stride, float* lookup_probs, int color, curandState*  
states) {  
    // Calculate global coordinates in the lattice (accounting for +1 padding)  
    int thread_x = blockIdx.x * blockDim.x + threadIdx.x + 1;  
    int thread_y = blockIdx.y * blockDim.y + threadIdx.y + 1;  
  
    // Boundary check: ensure the thread is within the LxL lattice  
    if (thread_x <= L && thread_y <= L) {  
  
        // Checkerboard update logic  
        // Only threads matching the current type are updated during this pass  
        // Here we used color because in general there is the map:  
        // even --> red | odd --> black  
  
        if ((thread_x + thread_y) % 2 == color) {  
            int global_index = thread_y * row_stride + thread_x;  
            int current_spin = lattice[global_index];  
  
            /* direct Global Memory Access */  
            int sum_neighbors = lattice[(thread_y - 1) * row_stride + thread_x] + // Top neighbor  
                               lattice[(thread_y + 1) * row_stride + thread_x] + // Bottom neighbor  
                               lattice[thread_y * row_stride + (thread_x - 1)] + // Left neighbor  
                               lattice[thread_y * row_stride + (thread_x + 1)]; // Right neighbor  
  
            // Map current state and neighbor sum to the 1D lookup table index  
            int row = (current_spin == -1) ? 0 : 1;  
            int col = (sum_neighbors + 4) / 2;  
  
            // Generate a random float using the pre-initialized CURAND state  
            float rand_val = curand_uniform(&states[global_index]);  
  
            // Metropolis Acceptance Criterion  
            if (rand_val < lookup_probs[row * 5 + col]) {  
                lattice[global_index] = -current_spin;  
            }  
        }  
    }  
}
```

CUDA - Shared memory kernel (1)



```
● ● ●

template <int BLOCK_SIZE>
__global__ void ising_step_shared(int* lattice, int L, int row_stride, float* lookup_probs, int color, curandState* states) {

    // local coordinates
    int thread_x = threadIdx.x;
    int thread_y = threadIdx.y;

    // global coordinates
    int x = (blockIdx.x * blockDim.x) + thread_x + 1;
    int y = (blockIdx.y * blockDim.y) + thread_y + 1;

    int global_index = y * row_stride + x;

    // shared memory
    __shared__ int my_cache[BLOCK_SIZE][BLOCK_SIZE];

    // load from Global to Shared
    if (x <= L && y <= L) {
        my_cache[thread_y][thread_x] = lattice[global_index];
    }

    // wait for all threads to load their spin
    __syncthreads();

    //...
}
```

1. Define local (block-relative) and global (lattice-relative) indices
2. Instantiate `my_cache`. The size is determined at compile-time via the `BLOCK_SIZE` template parameter
3. Synchronously load the lattice portion from Global to Shared memory using `__syncthreads()`

CUDA - Shared memory kernel (2)



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

```
● ● ●  
  
// UPDATE LOGIC  
if (x <= L && y <= L) {  
    if ((x + y) % 2 == color) {  
  
        int current_spin = my_cache[thread_y][thread_x];  
        int sum_n = 0;  
  
        // left neighbor  
        if (thread_x > 0) {  
            sum_n += my_cache[thread_y][thread_x - 1];  
        }  
        else{  
            sum_n += __ldg(&lattice[global_index - 1]);  
        }  
  
        // right neighbor  
        if (thread_x < BLOCK_SIZE - 1){  
            sum_n += my_cache[thread_y][thread_x + 1];  
        }  
        else {  
            sum_n += __ldg(&lattice[global_index + 1]);  
        }  
  
        // up neighbor  
        if (thread_y > 0){  
            sum_n += my_cache[thread_y - 1][thread_x];  
        }  
        else{  
            sum_n += __ldg(&lattice[global_index - row_stride]);  
        }  
  
        // down neighbor  
        if (thread_y < BLOCK_SIZE - 1){  
            sum_n += my_cache[thread_y + 1][thread_x];  
        }  
        else {  
            sum_n += __ldg(&lattice[global_index + row_stride]);  
        }  
    }  
}
```

1. Checkerboard algorithm based on "color"
2. Sum of the neighbors: if it belongs to the bulk → Shared Memory, otherwise read it from the Texture cache.

`__ldg()` utility: the hardware-managed Read-Only cache provides high-bandwidth access to "halo" spins that reside in Global Memory but are constant for the duration of the kernel.

CUDA - Shared memory kernel (3)



```
... ● ● ●

// Metropolis update by the means of the lookup table
// (allocated in the device within the constructor)
int lookup_index_0 = (current_spin == -1) ? 0 : 1;
int lookup_index_1 = (sum_n + 4) / 2;

// Generate a random float (0, 1] using the thread-specific state
// Each thread has a unique sequence offset to ensure statistical independence
// and eliminate race conditions on the generator state.
float rand_val = curand_uniform(&states[global_index]);

// Metropolis condition
if (rand_val < lookup_probs[lookup_index_0 * 5 + lookup_index_1]) {
    lattice[global_index] = -current_spin;
}
}
}
}
```

- Identify the probability in the lookup table and flip the spin accordingly

Steps CUDA



```
● ● ●

/***** UPDATE STEP CUDA *****/
void IsingModel2d::step_cuda_global(){
    // synchronize the padding
    launch_sync_padding_gpu(d_lattice, L, row_stride);
    // update even indices
    launch_ising_global(d_lattice, L, row_stride, d_lookup_probs, 0, d_states, cuda_block_size);

    launch_sync_padding_gpu(d_lattice, L, row_stride);
    // update odd indices
    launch_ising_global(d_lattice, L, row_stride, d_lookup_probs, 1, d_states, cuda_block_size);
}

void IsingModel2d::step_cuda_shared(){
    // synchronize the padding
    launch_sync_padding_gpu(d_lattice, L, row_stride);
    // update even indices
    launch_ising_shared(d_lattice, L, row_stride, d_lookup_probs, 0, d_states, cuda_block_size);

    launch_sync_padding_gpu(d_lattice, L, row_stride);
    // update odd indices
    launch_ising_shared(d_lattice, L, row_stride, d_lookup_probs, 1, d_states, cuda_block_size);
}
```

Synchronize the padding layers after each half-update



Table of Contents

1. Introduction

- Physics
- Hardware

2. Implementation

- General structure
- Serial implementation
- Parallel implementations

3. Benchmarks

- Physics correctness
- Performances

Physics analysis

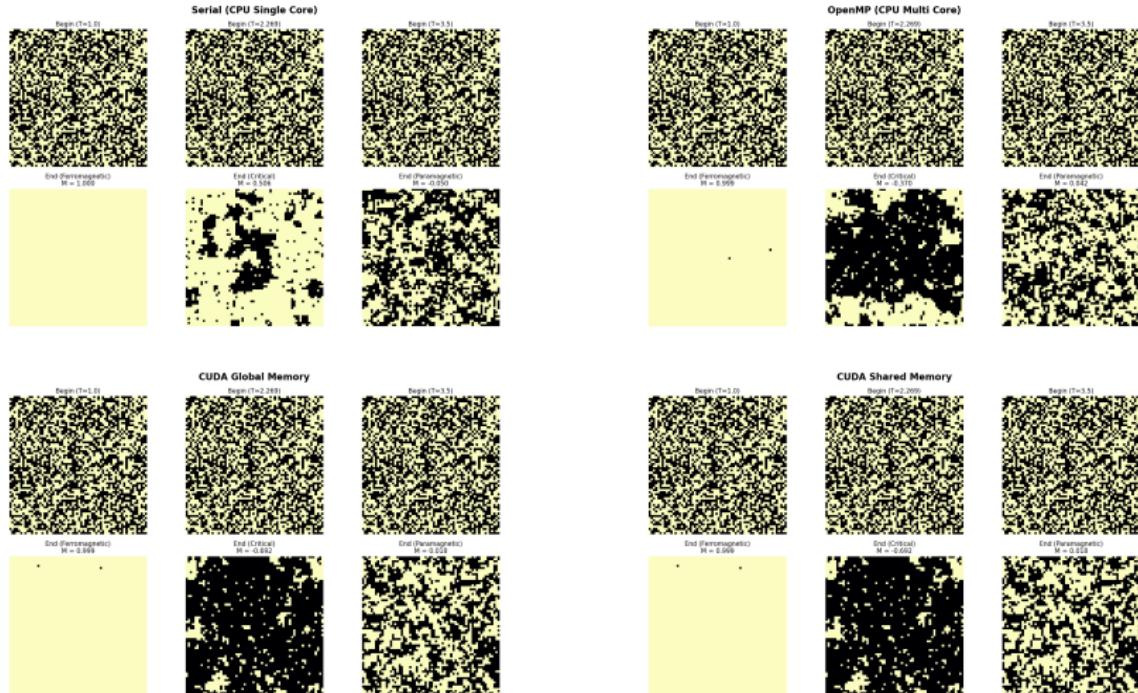


UNIVERSITÀ
DEGLI STUDI
DI PADOVA

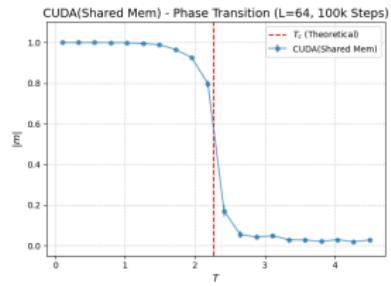
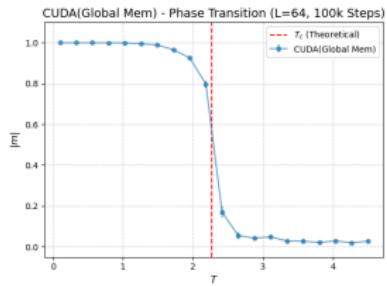
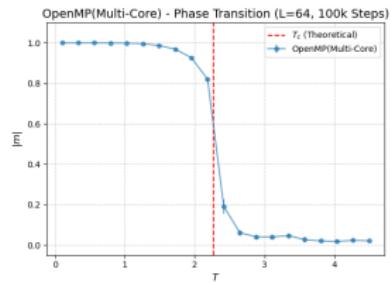
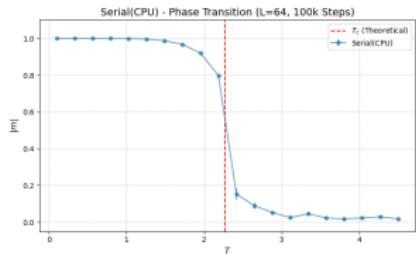
To check the algorithm correctness we will do the following tests:

- Spontaneous magnetization ($h = 0$)
- Energy minimization during the simulation steps ($h = 0$)
- (h, T) phase diagram

Visualizations

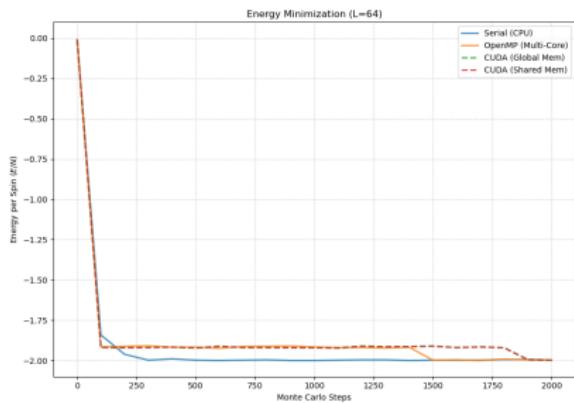


Phase transition $|m|$ vs T



Note: The deviation from the Onsager exact solution is caused by finite size effects ($L = 64$).

Energy minimization

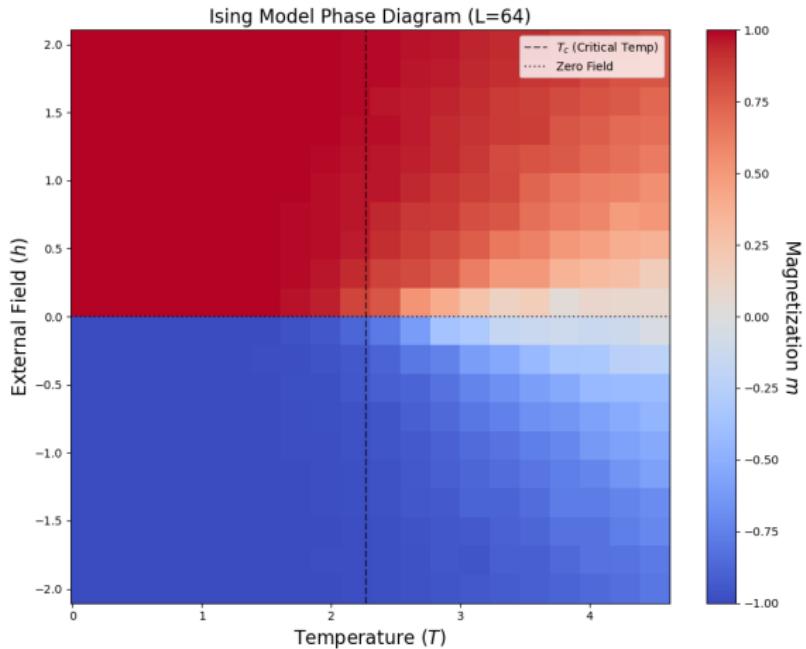


- All implementations successfully lead to the minimization of the system's energy.
- **Note:** The CUDAs implementations likely got trapped in a **metastable state** (local minimum).

Phase diagram (h, T)



UNIVERSITÀ
DEGLI STUDI
DI PADOVA



- Critical temperature T_c raises in presence of a magnetic field

Performance analysis



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

We check 4 different aspects:

- OpenMP scalability
- CUDA Kernel tuning
- Execution time and throughput analysis
- Limiting scaling

Benchmark Methodology

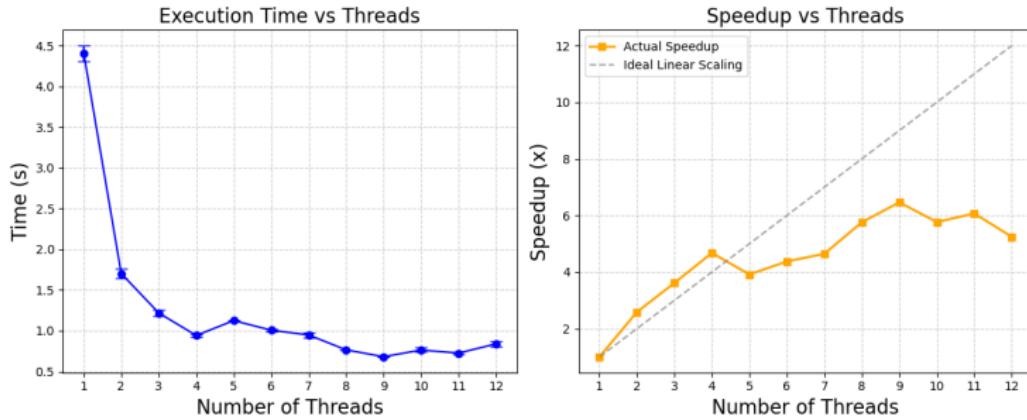


Time measurements were performed using the `time` module in Python, specifically the function `time.perf_counter()`:

- The measurement captures the total elapsed time on the Host (CPU).
- **CUDA Synchronization:** For GPU implementations, an explicit `cudaDeviceSynchronize()` is called before stopping the timer.
 - Without this, the timer would only measure the (negligible) kernel launch latency, as CUDA kernels are asynchronous.
- **Overheads:** Includes Python interpreter overhead and driver/kernel launch latencies (negligible for large L and high number of steps).

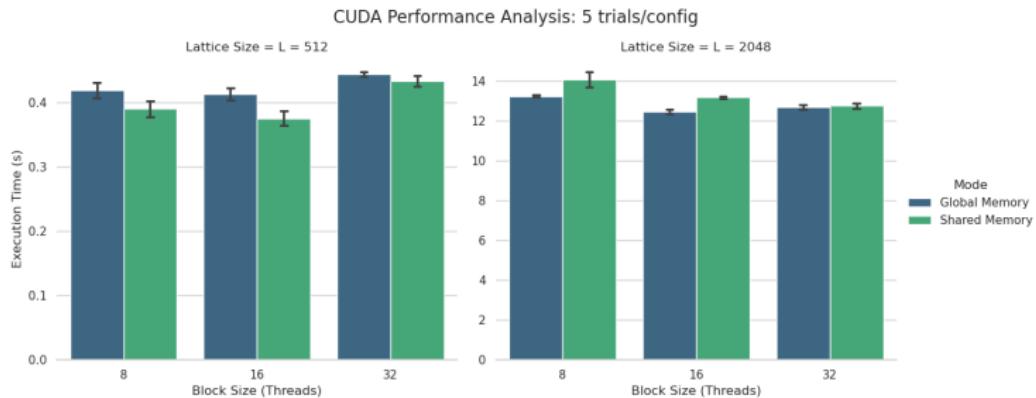
To be more precise and exclude overheads (on GPU), cuda events should be used instead

OpenMP vs n_threads



- In accordance with the CPU specifications (4 P-cores / 4 E-cores / 12 logical threads).
- Decrease at $n_threads = 12 \rightarrow$ (likely) context switching due to resource saturation.

CUDA vs block_size (1)



- **Latency-bound regime ($L \leq 512$):**
The problem is accessing the single data. The dataset fits within the L2 cache. **Shared Memory** outperforms Global Memory by minimizing access latency.
- **Bandwidth-bound regime ($L \geq 512$):**
The data size exceeds L2 capacity. Performance gap vanishes → simulation becomes limited by **VRAM bandwidth** rather than latency (see "Scaling" slide for better details).

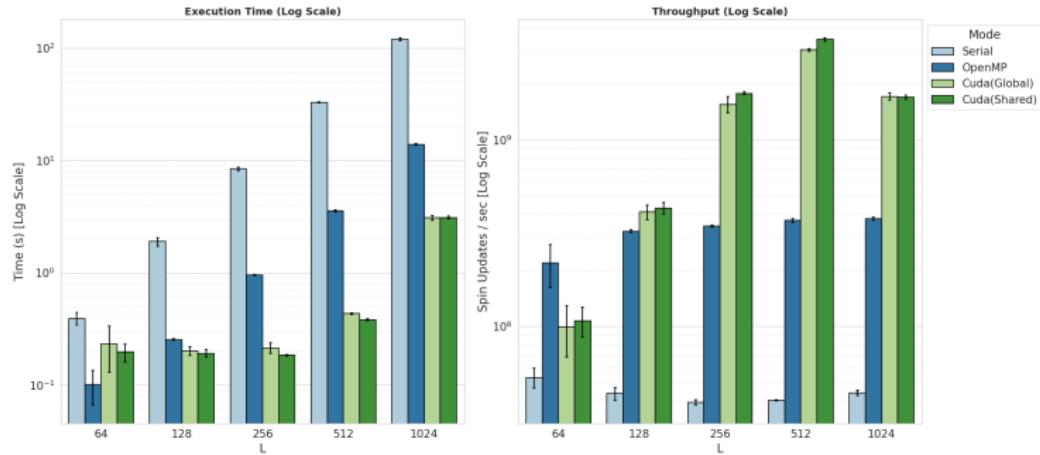
Performance analysis: block size effects



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

- **Small Blocks (8×8) at $L = 2048$:**
 - **Too many borders:** We load more "ghost cells" (neighbors) relative to the actual spins we update.
 - **Management Cost:** The GPU wastes time managing thousands of tiny tasks instead of calculating.
- **Large Blocks (32×32) at $L = 512$:**
 - **Wasted Space:** The blocks are too big to fit perfectly in the hardware, leaving empty slots.
 - **Waiting Time:** At `__syncthreads()`, if one thread is slow, 1023 others must wait doing nothing.

Execution time



- OpenMP is the fastest for small lattices (**less overhead costs**)
- GPU's results are in accordance with the previous block analysis



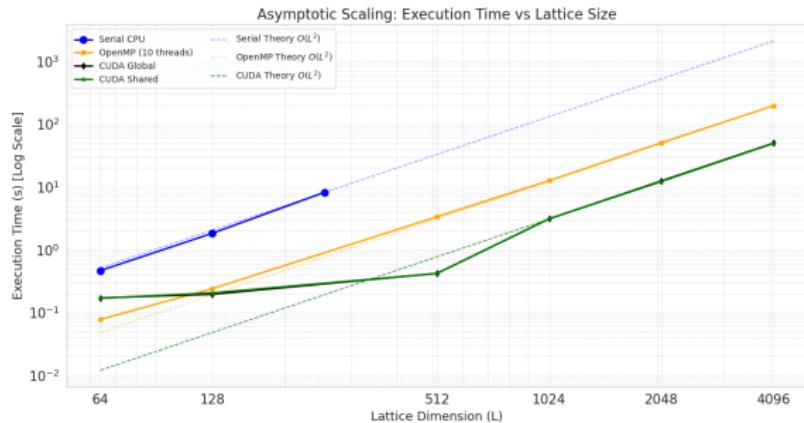
Throughput Analysis

In this analysis, the throughput is defined as the number of spin updates per second:

$$\nu_{\text{updates}} = \frac{L^2 \cdot \text{steps}}{T_{\text{execution}}}$$

- **CPU Implementation:** Throughput remains almost constant (or slightly decreasing) as L increases.
- **GPU Implementation:**
 - *Latency-bound Regime:* Throughput increases with L as hardware occupancy improves.
 - *Bandwidth-bound Regime:* Throughput reaches a plateau (becomes constant) as VRAM bandwidth saturates.

Asymptotic Scaling Analysis



- **CPU (Serial & OpenMP):** Both implementations exhibit the expected theoretical scaling of $\mathcal{O}(L^2)$ across the entire range of lattice sizes.
- **GPU (CUDA):**
 - **Small L :** dominated by constant overheads (kernel launch latency).
 - **$L \geq 512$:** the scaling converges to the asymptotic $\mathcal{O}(L^2)$ regime.



Conclusions & Future Work

- OpenMP is superior (or at least comparable) for small lattices ($L < 128$) where the overhead of PCIe data transfer and kernel launch outweighs the parallel benefits.
- For large systems ($L \geq 512$), the GPU implementation (Shared/Global) provides massive speedups, effectively hiding memory latency.
- The large L2 Cache on Ada Lovelace minimized the performance gap between Global and Shared Memory implementations.

Future Extensions:

- Scaling the domain decomposition using MPI to simulate massive lattices
- Implementing parallel reduction (e.g., Thrust) to calculate m and E directly on the GPU

Thank you for the attention

Random Number Generation - serial



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

- Private member `serial_rng` handles the spin flip probabilities
- It is an object of the standard library `<random>`, which:
 - Uses the **Mersenne Twister** engine (`std::mt19937`) for high-quality pseudo-random numbers
 - Provides a uniform distribution in the range $[0, 1)$ to be compared with the Metropolis acceptance probability P
 - Allows for **reproducible results** by using a fixed seed during the development and benchmarking phase
- **Note:** This generator is used for the Serial and OpenMP versions, whereas CUDA requires a different approach (`cuRAND`) to handle parallel generation

OpenMP - RNGs



- Naive parallel RNG (shared state) → **race conditions** or **serialization** (bottleneck)
- Solution: 1 RNG per thread (initialized in the constructor)

```
● ● ●  
  
// ... Constructor  
  
// initialize RNG engines for parallel threads in OPENMP  
int max_threads = omp_get_max_threads();  
for (int i = 0; i < max_threads; i++) {  
    omp_rngs.emplace_back(seed + i + 1);  
}  
// ...
```

CUDA - Helper functions (1)



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

Wrappers

```
● ● ●

void* gpu_alloc(size_t size) {
    void* ptr = nullptr;
    cudaMalloc(&ptr, size);
    return ptr;
}

void gpu_free(void* ptr) {
    cudaFree(ptr);
}

void gpu_memcpy_to_device(void* dest, const void* src, size_t size) {
    cudaMemcpy(dest, src, size, cudaMemcpyHostToDevice);
}

void gpu_memcpy_to_host(void* dest, const void* src, size_t size) {
    cudaMemcpy(dest, src, size, cudaMemcpyDeviceToHost);
}

void launch_cuda_sync(){
    /*Wrapper to synchronize blocks*/
    cudaDeviceSynchronize();
}
```

Class functions

```
● ● ●

// allocate memory on device
void IsingModel2d::allocate_cuda() {
    size_t lattice_bytes = row_stride * row_stride * sizeof(int);
    size_t states_bytes = row_stride * row_stride * 70; // excess CuRand state dimension estimation
    size_t lookup_probs = row_stride * row_stride * 10 * sizeof(float));
}

void IsingModel2d::copy_to_device() {
    // it copies the lattice from host to device
    gpuMemcpy_to_device(d_lattice, lattice.data(), lattice.size() * sizeof(int));
}

void IsingModel2d::copy_to_host() {
    // it copies the lattice from device to host
    gpuMemcpy_to_host(lattice.data(), d_lattice, lattice.size() * sizeof(int));
}

void IsingModel2d::deallocate_cuda() {
    // free the memory
    gpuFree(d_lattice);
    gpuFree(d_states);
    gpuFree(d_lookup_probs);
}

void IsingModel2d::upload_lookup_probs(){
    // upload to device the lookup probability table
    float host_lookup[10];
    for (int i = 0; i < 2; i++) {
        for (int j = 0; j < 5; j++) {
            host_lookup[(i*5 + j)] = lookup_probs[i][j];
        }
    }

    gpuMemcpy_to_device(d_lookup_probs, host_lookup, sizeof(float) * 10);
}

void IsingModel2d::device_synchronize(){
    launch_cuda_sync();
}
```

CUDA - Helper functions (2)



Synchronize the padding on the device

```
•••  
***** Sync padding *****  
__global__ void sync_padding_kernel(int* lattice, int L, int row_stride) {  
    int k = blockIdx.x * blockDim.x + threadIdx.x + 1;  
    if (k <= L) {  
        // Copy top/bottom rows  
        lattice[0 * row_stride + k] = lattice[L * row_stride + k];  
        lattice[(L + 1) * row_stride + k] = lattice[1 * row_stride + k];  
        // Copy left/right columns  
        lattice[k * row_stride + 0] = lattice[k * row_stride + L];  
        lattice[k * row_stride + (L + 1)] = lattice[k * row_stride + 1];  
    }  
}
```

```
•••  
void launch_sync_padding_gpu(int* d_lattice, int L, int row_stride) {  
    // use fixed number of threads for padding  
    int threads = 256;  
    int grid = (L + threads - 1) / threads;  
  
    sync_padding_kernel<<<grid, threads>>>(d_lattice, L, row_stride);  
}
```

- Used only one index to access the interested elements

- 1D block to launch the kernel (CUDA interprets both threads and grid as a dim3 object, even if are declared as int)

CUDA - Random Number Generation



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

We implemented a dedicated kernel and a host wrapper to efficiently handle the parallel RNG initialization using the **cuRAND** library.

Initialization Kernel

```
#include <cuda_runtime.h>
#include <curand.h>
#include <curand_kernel.h>
#include <curand_state.h>
#include <vector>
#include <cmath>
#include <assert.h>
#include <string.h>
#include <Ising_gpu_interface.h>
// **** Setup for RNGs ****
__global__ void curand_init(curandState* states, unsigned int seed, int L, int row_stride) {
    int x = blockIdx.x * blockDim.x + threadIdx.x;
    int y = blockIdx.y * blockDim.y + threadIdx.y - 1;
    if (x < L && y >= 0) {
        int index = y * row_stride + x;
        // initialize each thread with a Curandstate:
        // - seed: global starting point
        // - row_stride: sequence number to ensure independent, non-overlapping random sequences
        // - 0: initial sequence offset
        // - states[index]: destination in global memory for this thread's state
        curand_init(seed, index, 0, &states[index]);
    }
}
```

Host Wrapper

```
=====
// WRAPPERS (KERNEL LAUNCH)
=====

void launch_setup_rng(void* d_states, unsigned int seed, int L, int row_stride) {
    curandState* states = (curandState*)d_states;
    // fixed block size here
    dim3 block(16, 16);
    dim3 grid((L + block.x - 1) / block.x, (L + block.y - 1) / block.y);
    setup_rand_kernel<<<grid, block>>>(states, seed, L, row_stride);
    cudaDeviceSynchronize();
}
```

The `curandState` structure holds the generator's internal state (seed, sequence, offset) for each thread.

- `curand_init` is computationally expensive → execute it **only once** at the start.
- The states are stored in Global Memory and reused by the Metropolis kernels at each simulation step.

CUDA - Global memory kernel (wrapper)



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

```
● ● ●

void launch_ising_global(int* d_lattice, int L, int row_stride, float* d_lookup_probs, int color, void* d_states,
int block_size){
    /*Wrapper to launch 'global' algorithm*/

    // cast to the correct type (so g++ does not raise errors)
    curandState* states = (curandState*) d_states;
    if (block_size > 32){
        std::cout << "It is not possible to have a block size of "<< block_size << " --> casted to 32" << std::endl;
        block_size = 32;
    }

    dim3 block(block_size,block_size);
    dim3 grid((L + block.x - 1)/block.x, (L + block.y - 1)/block.y);

    ising_step_global<<<grid,block>>>(d_lattice, L, row_stride, d_lookup_probs, color, states);
}
```

- Control condition regarding the limit of 1024 threads in a single block.

CUDA - Shared memory wrapper



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

```
● ● ●

void launch_ising_shared(int* d_lattice, int L, int row_stride, float* d_lookup_probs, int color, void* d_states,
int block_size) {
    // cast to the correct type (so g++ does not raise errors)
    curandState* states = (curandState*)d_states;

    dim3 block(block_size, block_size);
    dim3 grid((L + block.x - 1) / block.x, (L + block.y - 1) / block.y);

    // Choose the correct template based on runtime block_size
    switch(block_size) {
        case 8:
            ising_step_shared<8><<<grid, block>>>(d_lattice, L, row_stride, d_lookup_probs, color, states);
            break;
        case 16:
            ising_step_shared<16><<<grid, block>>>(d_lattice, L, row_stride, d_lookup_probs, color, states);
            break;
        case 32:
            ising_step_shared<32><<<grid, block>>>(d_lattice, L, row_stride, d_lookup_probs, color, states);
            break;
        default:
            std::cerr << "ERROR: Block size " << block_size << " not supported! (Use 8, 16, or 32)" << std::endl;
    }
}
```

- Switch to define block size (both for blocks and shared memory)

Update method



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

Serial

```
***** UPDATE STEP *****

void IsingModel2d::update(Mode mode, int steps) {
    switch(mode) {

        //***** SERIAL CPU EXECUTION *****
        case Mode::serial:
            for (int i = 0; i < steps; i++) {
                // The serial step already includes sync_padding() internally
                step_serial();
            }
            break;
    }
}
```

OpenMP

```
//...
//***** PARALLEL OPENMP EXECUTION ****/
case Mode::openMP:
    step_openmp(steps);
    break;
//...
```

CUDA

```
***** PARALLEL CUDA EXECUTION ****/
case Mode::cuda_global:
    // d_lattice must be already allocated and populated on the GPU.
    // If this is the very first run, ensure copy_to_device() was called before.
    for (int i = 0; i < steps; i++) {
        // perform one full Monte Carlo step (Padding + Red + Black)
        step_cuda_global();
    }

    // wrapper for the blocks synchronization
    launch_cuda_sync();

    // note: We do NOT copy data back to Host here.
    break;

case Mode::cuda_sharing:
    // d_lattice must be already allocated and populated on the GPU.
    // If this is the very first run, ensure copy_to_device() was called before.
    for (int i = 0; i < steps; i++) {
        // perform one full Monte Carlo step (Padding + Red + Black)
        step_cuda_sharing();
    }

    // wrapper for the blocks synchronization
    launch_cuda_sync();

    // note: We do NOT copy data back to Host here.
    break;
}
```

Speedup for L = 512 (5k steps)



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

Mode	Time (s)	Speedup (vs Serial)
Serial	32.7749	1.00x
OpenMP	3.5483	9.24x
CUDA (Global)	0.4315	75.96x
CUDA (Shared)	0.3793	86.42x

Extra: Simulated annealing (1)



Simulated Annealing is an optimization method that adapts the system's temperature according to a specific cooling schedule during the simulation.

The decay formula is the following:

$$T(t) = T_0 \cdot \exp(-\alpha t)$$

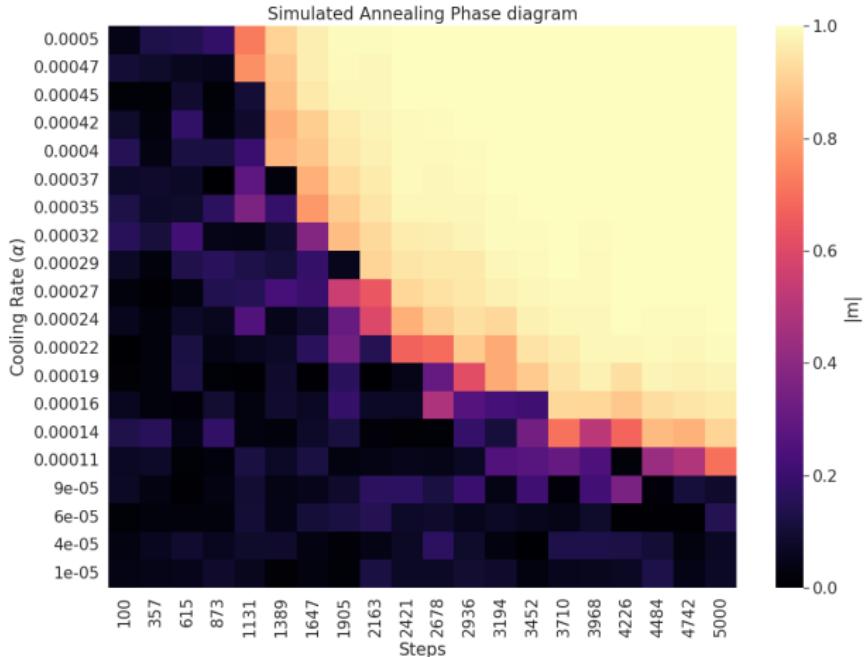
- T_0 : Initial temperature
- α : Decay rate (cooling speed)
- t : Simulation step

Extra: Simulated Annealing (2)



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

Parameter Space Analysis (α vs Steps) (Starting temperature $T_0 = 4$)



Extra: Simulated Annealing (3)



UNIVERSITÀ
DEGLI STUDI
DI PADOVA

- **Low α (Slow cooling):** The temperature decays too slowly. Within the limited steps, the system remains at high T and remains disordered due to **thermal fluctuations**.
- **Increasing α :** The cooling rate is sufficient to reach the ordered phase (Ground State), allowing the system to **escape local minima** and settle before freezing.

